



Faculty of Engineering and Architecture

Building a better VHDL testing environment

Joren Guillaume

Supervisors: prof. ir. L. Colman, dr. ir. H. Eeckhaut
Counsellor: ir. ing. L. Lemiengre

Master's dissertation submitted in order to obtain the academic degree of
Master of Science in Electronics and ICT Engineering Technology

Academic year 2014–2015



Faculty of Engineering and Architecture

Building a better VHDL testing environment

Joren Guillaume

Supervisors: prof. ir. L. Colman, dr. ir. H. Eeckhaut
Counsellor: ir. ing. L. Lemiengre

Master's dissertation submitted in order to obtain the academic degree of
Master of Science in Electronics and ICT Engineering Technology

Academic year 2014–2015

Usage restrictions

The author gives permission to make this master dissertation available for consultation and to copy parts of this master dissertation for personal use. In the case of any other use, the copyright terms have to be respected, in particular with regard to the obligation to state expressly the source when quoting results from this master's dissertation.

Preface

In the past year and a half I had the pleasure to work on the concept of integrating the software and hardware development worlds. In this thesis is contained the culminated knowledge and experience that was obtained during this period.

In the first part the tackled problem is explained, along with a summarised background on the practices and concepts one needs to know to fully understand the next part. In the second part the practical side of the work is demonstrated and thoughts are shared on the usability and future of this line of work.

Acknowledgements

In the course of learning, producing and revising both the code and the written text many helping hands guided my way. With this I would like to extend my gratitude to them.

To my supervisors prof. ir. Luc Colman and dr. ir. Hendrik Eeckhaut. The former for providing his many years of experience in supporting students, keeping a calm but firm grip on matters and always pushing to the next step. The latter for granting me the subject, providing his technical expertise and willingness to work on my problems whenever I asked.

Most certainly to my counsellor ir. ing. Lieven Lemiengre for providing an endless amount of feedback and brainstorming whenever. For providing technical expertise and insight when I had none to be found, for providing ideas and angles to keep me going. For remaining realistic and setting achievable goals when the light at the end of the tunnel was darkest. And of course, for reviewing and critiquing this thesis.

To my fellow students for the pleasant atmosphere that dominated most of my studies. For providing help when asked and receiving advice without reluctance. And for those long hours spent in the telecommunications lab working side by side on many things.

Finally, to my darling fiancée Dorine Ndagsi for supporting me the many years we have been together. For providing a critical eye when needed and bringing me down a notch when explanations were overly complicated. And most certainly for being the driving force behind most of the work that was done at home, at work and in life.

To all those that are not named, you are not forgotten, the list would simply be endless.

Abstract

English

Digital electronics is a global industry of undeniable importance. Producing these electronics is usually done with the aid of computer programs analogous to software development. In this thesis a number of practices from the software world are transferred to the hardware world. It is investigated whether unit testing has merit in hardware development and whether test-driven development is applicable. Furthermore, continuous integration as a method of speeding up and automating development with software development standards is tested. These practices are brought together in the form of a Python script that uses ModelSim to compile and simulate the VHDL source code and test benches.

Keywords: VHDL; TDD; unit testing; continuous integration; test benches

Nederlands

Digitale elektronica is een wereldwijde industrie van vitaal belang. Deze elektronica wordt gebruikelijk geproduceerd met behulp van computerprogramma's zoals bij software ontwikkeling. In deze thesis worden een aantal gebruiken van de software wereld overgezet naar de hardware wereld. Er wordt onderzocht of unit testing van nut kan zijn in de hardware ontwikkelingswereld, en of test-driven development toepasbaar is. Verder wordt continuous integration getest als een manier om sneller en automatisch met software ontwikkelingsstandaarden te ontwikkelen. Deze praktijken worden samengebracht door een script, geschreven in Python, dat ModelSim gebruikt om de VHDL broncode en test benches te compileren.

Trefwoorden: VHDL; TDD; unit testing; continuous integration; test benches

Building a better VHDL testing environment

Joren Guillaume

Supervisors: prof. ir. Luc Colman, dr. ir. Hendrik Eeckhaut, ir. ing. Lieven Lemiengre
Ghent University

Abstract— The VHSIC Hardware Description Language (VHDL) is a language used industry-wide for developing digital electronics. This paper explores the possibility of using software development practices in hardware development that uses VHDL. Borrowing elements from test-driven development, specifically unit testing, a test bench is split into parts that are executed sequentially. To achieve this, a script is written in Python that uses the ModelSim compiler/simulator. This script is then automatically called by a continuous integration server that captures and processes the results.

Index terms—VHDL, TDD, unit testing, continuous integration, test benches

I. INTRODUCTION

Developing digital hardware is and will remain a global industry of undeniable importance. VHDL is one of the important hardware design languages, having been used in this industry for many decades. Although improvements have been made to the language itself, the development practices have been lagging behind the software world for several years. As with all production, any improvements that can be done to the development process with appropriate cost should be a welcome sight. However, the industry is slow to adapt or convert to more up-to-date processes. This paper aims to investigate a number of software development practices and build a working framework in which they can be used for hardware design [1–3].

A. Unit testing

In a testing environment, a unit is the most basic behaviour of a piece of code that can be tested. It is called a unit because there is nothing smaller. Unit testing is the practice of taking these small parts and writing tests for them and only them, so that all tests are performed on units. A test only performs on a single unit of code; if the test uses parts of code outside of the unit, it is not a unit test. One of the biggest advantages of this testing method is the precise debugging information that is given should a test fail. After the pinpointing, only a small part of the code needs to be examined, making it very efficient [4, 5].

B. Continuous integration

Continuous Integration (CI) is the practice of automatically and very rapidly integrating code updates into a build. It encompasses a number of practices such as:

- Revision control
- Automated building
- Automated testing

CI aims to prevent the trouble in development that arises when developers wait too long to integrate their edits into the main build. It promotes early and rapid integration by automating these steps and downloading all the latest code from a repository. This is assuming that the repository contains the newest code and is updated on a frequent basis. After integrating and building the code, it should run the tests on this code and read the results. If all of these steps are automated and triggered on a timed basis, developers should have a steady stream of progress reports from the test results available [6, 7].

II. METHODS

Transferring the concepts from software to hardware required a number of concessions that would otherwise make implementation impossible.

A. Unit testing in VHDL

A unit test was defined as a test or group of tests that could be run at earliest convenience. This sometimes needed inputs to be set and the simulation to be run up until the required point. All test groups were identified manually with markers in the original test bench code.

B. Python script

A script was written in Python to combine all of the aspects investigated. This script is called from the command-line or a shell and takes a number of arguments to determine method of separation, source test benches, output location and VHDL version. It read the test benches and separated them into new test benches. These test benches were then compiled and simulated using ModelSim and condensed reports of the output were made in the JUnit XML specification.

C. Bitvis utility library

A utility library made by Bitvis was used to speed up development. Functions in this library made it possible to quickly produce tests with uniform output. The library's default output was slightly modified to better integrate with the Python script's report generation.

D. Hudson-CI

Continuous integration was implemented using the Hudson-CI program. A job was made to download the code from an on-line Git repository. The job then executed the Python script and captured its reports. Hudson provided built-in support for JUnit XML report reading and automatically provided statistics on failure and success.

III. RESULTS

Separating the tests at earliest convenience provided the benefit of not halting tests when a critical error occurred. Instead, only that single group of tests would return as failed to run. The CI server gave clear indications which tests had passed and failed, but was not able to correctly read the nano- and picosecond timing a hardware simulation requires. Furthermore, the script was limited when it came to the complexity of the test benches it could read, putting boundaries on which projects could be tested.

IV. CONCLUSION

It was proven that software techniques could be modified to work with a hardware development environment, but that more detailed and thorough work is required before it reaches full potential.

REFERENCES

- [1] Roger Lipsett, Erich Marschner, and Moe Shahdad. "VHDL - The Language". In: *IEEE Design & Test of Computers* 3.2 (Apr. 1986), pp. 28–41. doi: 10.1109/mdt.1986.294900.
- [2] R.D. Acosta, S.P. Smith, and J. Larson. "Mixed-mode simulation of compiled VHDL programs". In: *1989 IEEE International Conference on Computer-Aided Design. Digest of Technical Papers*. IEEE Comput. Soc. Press, 1989. doi: 10.1109/iccad.1989.76930.
- [3] T.E. Dillinger et al. "A logic synthesis system for VHDL design descriptions". In: *1989 IEEE International Conference on Computer-Aided Design. Digest of Technical Papers*. IEEE Comput. Soc. Press, 1989. doi: 10.1109/iccad.1989.76906.
- [4] Don Wells. *Code the Unit Test First*. 2000. URL: <http://www.extremeprogramming.org/rules/testfirst.html> (visited on 03/20/2014).
- [5] Martin Fowler. *UnitTest*. 2014. URL: <http://martinfowler.com/bliki/UnitTest.html> (visited on 06/30/2014).
- [6] Fazreil Amreen Abdul and Mensely Cheah Siow Fhang. "Implementing Continuous Integration towards rapid application development". In: *2012 International Conference on Innovation Management and Technology Research*. IEEE, May 2012. doi: 10.1109/icimtr.2012.6236372.
- [7] Sean Stolberg. "Enabling Agile Testing through Continuous Integration". In: *2009 Agile Conference*. IEEE, Aug. 2009. doi: 10.1109/agile.2009.16.

Contents

I	Problem and background	1
1	Problem	1
2	Introduction	2
2.1	Hardware description languages	2
3	Current industry practices	4
3.1	Assertions	4
3.2	The generic test bench	4
3.3	Coverage	4
3.3.1	Code coverage	5
3.3.2	Functional coverage	5
3.4	Verification	5
3.4.1	Constrained random verification	5
3.4.2	Formal verification	5
3.5	Simulation tools	5
3.5.1	ModelSim	6
4	Software development practices	7
4.1	Test-driven development	7
4.1.1	Unit testing	7
4.1.2	Test first development	7
4.1.3	Refactoring	7
4.1.4	Test-driven development	8
4.2	Continuous integration	8
4.2.1	Revision control	8
4.2.2	Build automation	8
4.2.3	Test automation	9
4.3	xUnit	9
4.4	Python	9
II	Developing a framework	10
5	Outlining	10
5.1	First draft	10
5.2	Draft review	10
5.3	Fresh start	11
6	Using the framework	12
6.1	Preparing the test bench	12
6.2	Calling the parser	13
6.2.1	Arguments to the parser	13
6.2.2	Processing methods	14
6.3	Report and log generation	14

6.3.1	Basic report	14
6.3.2	XML report	14
6.3.3	Logkeeping	15
6.4	Hudson-CI	15
6.4.1	Using Hudson-CI	15
6.5	Bitvis utility library	16
7	Results	17
8	Future work	18
8.1	Wider tool support	18
8.2	Lexical analysis	18
8.3	Adapted CI tool	18
9	Commentaries on developing VHDL	19
9.1	An industry stuck in 1993	19
9.2	Lack of reflection or introspection	19
9.3	Hardware developers are not software developers	19
10	Conclusion	20
	References	21
	Appendices	24
	Appendix A Code examples	24
A.1	DFF test bench	24
A.2	Revised DFF test bench	25
	Appendix B Figures	26
B.1	Hudson-CI example statistics	26

List of Figures

1	Typical design flow, red indicates the main focus of this thesis.	3
2	Three step TDD design flow	8
3	Workflow of the framework	11
4	Hudson jobs status window	17
5	Hudson-CI example statistics.	26

List of Acronyms

AES	Advanced Encryption Standard	17
CC	Code Coverage	5
CI	Continuous Integration	8, 9, 15, 18
CRC	Cyclic Redundancy Check	17
CRV	Constrained Random Verification	5
DFF	D Flip-Flop	2, 12, 16
DT	Directed Testing	5
DUT	Device Under Test	3, 4
FC	Functional Coverage	5
FSM	Finite State Machine	5
HDL	Hardware Description Language	2, 4–6
IEEE	Institute of Electrical and Electronics Engineers	2
RC	Revision Control	8, 15, 20
RTL	Register Transfer Level	2
SHA	Secure Hash Algorithm	17
TDD	Test-Driven Development	7
TFD	Test First Development	7
UUT	Unit Under Test	3, 5
VHDL	VHSIC Hardware Description Language	1, 2, 6, 10, 16, 17, 19, 20
VHSIC	Very High Speed Integrated Circuits	1
XML	eXtensible Markup Language	9, 14, 16, 20

Part I

Problem and background

1 Problem

Developing digital hardware with the Very High Speed Integrated Circuits (VHSIC) Hardware Description Language (VHDL) is, like any code, prone to errors; either made by the developer or due to wrong product specifications. To ensure these errors are weeded out before the more expensive roll-out or production phase begins, the code must be subjected to rigorous testing. This testing process is long and extensive, and requires specially trained engineers.

Because testing is such a time-consuming process, finding errors often results in severe delays. This is due to the need to correct the error as well as the need to test for other errors. Therefore, it is in the best interests of both testing and development engineers to find and correct errors with minimal delay and maximal efficiency. Not to mention developing code with as little as possible errors from the start. It is also important to create test benches with maximal coverage, optimal readability and minimum time spent. This entire process should affect the least amount of code possible in order to minimize time spent retesting and modifying the code.

In this thesis, the objective is to explore the possibility of creating an operating system independent framework. This framework should allow developers to quickly and consistently create, execute and evaluate their code with specially modified test benches. To accomplish this, a number of industry standard tools will be incorporated around a central Python script. This framework should ensure easy test bench creation, timely and automated building, automated testing and readable test report generation.

2 Introduction

2.1 Hardware description languages

A Hardware Description Language (HDL) can be used to describe digital electronics, i.e. hardware, in different levels. The level that uses certain blocks of logic gates to describe more complex behaviour is called the Register Transfer Level (RTL). Some blocks are standard implementations that have been widely used and are nearly fully optimized, such as memories, flip-flops and clocks. One of the widely used HDLs is VHDL, having been developed in the 1980s. Originally it was meant to have a uniform description of hardware brought in by external vendors to the U.S. department of defence. It was quickly realized that simulation was possible with a correct description and tools evolved to be used as such [1, 2]. The final step was to create tools that could not only simulate, but also synthesize (i.e. create actual hardware layouts) from these descriptions [3–6]. An RTL flip-flop implementation written in VHDL is shown here:

```
LIBRARY IEEE;
USE IEEE.std_logic_1164.ALL;

ENTITY dff IS
    PORT(d      : IN  std_logic;
         clk     : IN  std_logic;
         q       : OUT std_logic;
    END dff;

ARCHITECTURE Behavioural OF dff IS
BEGIN
    PROCESS (clk)
    BEGIN
        IF rising_edge(clk) THEN
            q <= d;
        END IF;
    END PROCESS;
END Behavioural;
```

The Institute of Electrical and Electronics Engineers (IEEE) 1164 library provides a number of extensions on the original VHDL IEEE 1076 specification. These extensions allow a more realistic simulation and description of hardware behaviour [7]. An *entity* defines the inputs and outputs of a certain building block, in this case the D Flip-Flop (DFF). The *architecture*, in this case named *Behavioural*, takes the description of an entity and assigns a real implementation to it. In this architecture, we have one process and it is *sequential*, meaning that every update in the process follows an update of *clk*, a clock signal. The *d* stands for delay, and it simply puts on its output *q*, that which was on the input *d*, one clock period earlier. All processes are executed in parallel, which does not mean that all are triggered at the same time, nor that they take equally as long to finish. It means that any process can be executed alongside any other process. In this case there is only one (nameless) process that describes the entire behaviour of the flip-flop. The process waits for the rising edge of the clock, which is a transition from zero to one. Subsequently it schedules the value of *d* to be put on *q* until the next rising edge appears.

This is a basic example of an entity, an architecture and a process. Before the code can be put to use in a working environment, it first needs to be tested. This is done through the use of *test benches* [8]. Test benches are made up of code that takes a certain building block, the Unit Under

Test (UUT) or Device Under Test (DUT). The specialized tools mentioned in the previous paragraphs take the test benches and read them. Inside are usually instructions to put a certain sequence of values on the inputs and monitor the outputs [9]. Applying this to the example listed earlier, the test bench would contain a process with a clock, local signals coupled with the ones in the entity, some stimuli and *wait* statements. The signals are linked to the DFF, which is now the UUT. Then the clock starts ticking and the input *d* is made ‘0’ or ‘1’ every now and then. All that is left is to assure the output *q* is always ‘0’ and ‘1’ exactly one clock cycle later. The full code is listed in appendix A.1.

If the device performs normally, the received output sequence should match an expected output sequence. In these test benches it is good practice to also observe how well a device performs if its inputs behave outside of the normal mode of operation. When all of these tests have finished and the output performs as expected, the device is ready to be moved further down the developmental process [10].

The higher level components, such as a registry, employ a number of the lower level ones to create more complex logic [6]. The flip-flop could be used in certain numbers to build a *register*. A register is a collection of ones and zeroes (henceforth referred to as *bits*) that is used to (temporarily) store these values. The register could then be used alongside combinational logic to build an even bigger entity. The idea here is that small building blocks can be combined to produce vast and complex circuits that are nearly impossible to describe in one go [11]. Adding all these layers together also creates a lot of room for error, and having a multi-level design makes it challenging to pinpoint the exact level and location of any errors. If a device is not tested properly and faults propagate throughout the next stages of development, they can be very expensive to correct [12]. Therefore, a large portion of time is spent writing and executing tests, which provide feedback for the developmental cycle. In figure 1 an overview of the design flow is given up until synthesis.

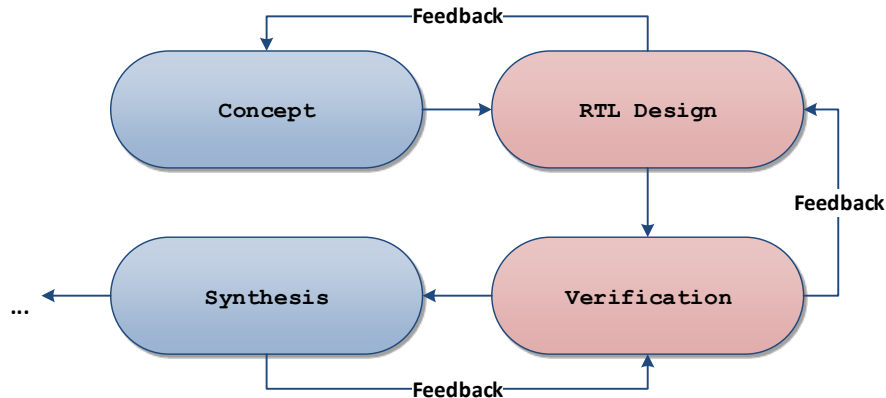


Figure 1: Typical design flow, red indicates the main focus of this thesis.

3 Current industry practices

A number of practices exist to improve speed and quality of testing and coding. Many of these practices are already applied in the software development industry, where this is sometimes quite literally their entire business. Moving to HDLs, it stands to reason that code meant for synthesis may not be able to follow all of these practices. However, the industry has formed several practices of its own to try and create a uniform verification process. The most well-known will be discussed in some detail below.

3.1 Assertions

Assertions are the standard practice of verification. An assert in VHDL is very simple: check whether some boolean condition is true [13]. If true, do nothing. If false, return the error message and throw an error of a certain level, as in the following example:

```
assert (not q) report "Unexpected output value" severity failure;
```

In this example, an error is thrown of the severity *failure*, which ends the simulation, if the value on the output *q* (see section 2.1)) isn't *logically* true. In this case it means the output has to have a value of *high* or *1*. The severity of the assertion can be in the range of notice to failure, and the simulator might be set to respond or stop only to a certain level of severity. Doing this for the right values at intervals gives the developer a quick overview of whether everything went according to plan. After all, if things went wrong, the assertion should have thrown an error here or there.

3.2 The generic test bench

The generic test bench operates as mentioned in section 2.1. It assigns to every input and output on the DUT their counterpart in the test bench and contains a number of processes with stimuli [14]. Standard procedure is:

1. Wait for some clock periods to 'ready' the design
2. Apply stimuli to the inputs
3. Wait for the appropriate number of clock cycles
4. Use asserts to check whether the outputs have the right values
5. Repeat steps 2 through 4 until satisfied
6. End with an infinite *wait* to suspend the process

Creating this kind of test bench for a large, hierarchical project would certainly become lengthy and unclear. To counter these disadvantages, a number of practices were created to maintain control over which designs are tested and how the testing is done. The most commonly used practices are explained below.

3.3 Coverage

Coverage is a generic term that is used to describe how fully a design has been tested on one aspect or another. There exist many tools for different coverage analyses, but in this section we will focus only on the types of coverage.

3.3.1 Code coverage

In development, Code Coverage (CC) is a type of measurement to indicate how well the source code has been tested. With the use of a coverage report, unused blocks of code can be found, which might indicate unnecessary code or a bug. Imagine a Finite State Machine (FSM) with an unused reset state, this might indicate that the reset isn't functioning properly or that there are no tests covering the reset. However, CC does not provide any real functional analysis. It does not indicate any missing lines of code nor does it tell the developer whether the inputs and outputs behave properly [15, 16].

3.3.2 Functional coverage

Functional Coverage (FC) is the practice of measuring whether the status of the UUT meets with certain specifications at specific times during the testing process. These specifications are created by the developer and are used to check whether the design performs as expected. Good practice is to include corner cases, cases that cover very rare occurrences and so on. That way, the device is sure to be in working condition even under unexpected circumstances. As such, FC provides the type of coverage CC lacks [17, 18].

3.4 Verification

3.4.1 Constrained random verification

Constrained Random Verification (CRV) is an industry practice where one or more inputs are generated randomly, within certain bounds or *constraints* [19, 20]. The output is then accepted as correct if it too is within certain (other) constraints. This practice was brought into use after designs grew too large for Directed Testing (DT) to support. DT has verification engineers write out very specific things they want to test. For instance, a reset pulse to verify that the reset is working correctly. CRV opposes this with the idea that for all behaviour to be tested properly in large designs, the amount of time spent writing and executing tests would simply become too great. It proposes a solution where inputs are generated randomly, within certain bounds, but in a sufficiently large quantity to have implicitly covered all scenarios. It is important to note that with DT expected behaviour is directly tested, but with CRV it is likely to be the unexpected behaviour that gets tested too. This solved the long-standing problem of testing any behaviour, including the unexpected.

3.4.2 Formal verification

On top of the aforementioned, there are several more practices involving unique ways of verifying the properties of a design, but that aren't used sufficiently to merit full detailing. Formal verification is such a practice, where the core idea is to mathematically prove the design from its specifications [13, 21]. The upside is that the design is completely verified. However, it has gone out of use, because proving large designs is not only tedious, but takes up large amounts of man-hours as well.

3.5 Simulation tools

As mentioned in section 2, HDLs are used for developing hardware, and need to be tested and built as such. Like many other programming languages, they need a dedicated compiler to fault-check the code and build the binaries. Unlike those other programming languages they also

need a simulator in order to verify the builds. There exist many compilers and simulators, almost none exclusive for VHDL. Nearly all of them are dual-language and support some form of Verilog, a different HDL. However, many of them refuse to support the latest additions to the VHDL language specification, even the 2002 additions are hard to be found [22–25].

3.5.1 ModelSim

ModelSim is the simulator that was investigated for several reasons. Firstly, a free student edition was available. Considering licenses outside of the educational system can easily cost upward of \$25,000 (€20,250 at time of writing) this made it ideal for any thesis or student-related work. Proof of this is the extensive use of ModelSim in the courses related to HDL development. Secondly, ModelSim supports all versions of VHDL, which is the HDL we are processing. Even in 2014, only one other tool was found to support all of VHDL-2008, which is Aldec’s *Active-HDL*. And last but not least, ModelSim is one of the industry’s most used simulators, enabling a pool of experience to be consulted [26].

4 Software development practices

During the writing of this thesis, a number of approaches were investigated and some were put to practical use. In this chapter, the more useful and tried of the aforementioned are discussed in detail.

4.1 Test-driven development

Test-Driven Development (TDD) is a proven development technique that has regained traction in the past decade, primarily through the efforts of Kent Beck [27, 28]. This practice has proven to increase test coverage, decrease defect density as well as improve code quality [29–31]. The technique focuses on tests being created before the actual code. It is important to make certain distinctions before going more in depth on the used methods. The developing community has a great many practices, each with their own names and methods, and almost none are mutually exclusive.

4.1.1 Unit testing

To understand TDD, a basic knowledge of *unit testing* is required. In software testing, a unit test is a test designed to check a unit of code [27, 32]. Ideally, this unit is the smallest piece in which the code can be divided. A unit test should always test only a single entity, and only one aspect of that entity’s behaviour. This division in units has a number of benefits, one of the most important being that code is exceedingly easy to maintain. Furthermore, the division of the code makes changes, when needed, fast to be carried out and ensures that only the modified code needs to re-tested.

4.1.2 Test first development

Another main component of TDD is Test First Development (TFD), a technique that has a developer writing tests before any code has been written [32]. This method puts behaviour as the primary drive for development. After all, if the test is conceived before the code it tests, the goal is to make the tests pass by implementing the behaviour it tests. This has developers thinking of what the code should do rather than what the exact implementation has to be. A key feature of a new test is that it has to fail during its first run. If not, the test is obsolete, as the functionality it tests has already been implemented. This is called the *red* step because of the colour it usually shows in the results. After being run for the first time (and failing), the developer implements just enough code to make the test pass; this is called the *green* step. Once all tests succeed, it is time for a new test.

4.1.3 Refactoring

The third pillar of TDD is refactoring. After the tests succeed, it is necessary to clean up the code [33]. A well-done TFD implementation uses the bare amount of code needed to make the test pass, but this might not be the best code possible. Refactoring means that you take existing code and modify it to perform better. This should leave both the input and output of the tests and code unchanged, only the way the code processes input is altered. During this step, it is important that no parts of the test code, outputs or inputs are changed. Otherwise, either the test would behave differently or new tests would have to be written.

4.1.4 Test-driven development

Test-Driven Development is a combination of the previously mentioned techniques. A unit test is written before any code is, the test is then executed and should fail (the red step). After the failure, the most basic code to make the test pass is implemented. The test is then executed again and should pass (the green step). After this first pass, the code written for the passing of the test, and only this code, is edited to perform better and be less complex. This follows all steps mentioned above, a *unit test* is written according to *test first development* and is *refactored* later. The classic flowchart is shown in figure 2.




Figure 2: Three step TDD design flow

4.2 Continuous integration

Continuous Integration (CI) is a software development technique in which developers upload the edits on the software, or their *build*, to a central server which then *continuously integrates* the code, usually from multiple developers. This is intended to prevent integration headaches when the code of several simultaneous changes has diverged too greatly. Usually the code has diverged to such an extent that it would take much more time to make the edits work together than when they had been integrated early on.

4.2.1 Revision control

There are many aspects to a properly maintained CI system, but one key aspect of overall programming has to be Revision Control (RC). Not to be confused with the "undo" button in the preferred editor, a good RC system does allow for any and all mistakes made over different edits to be undone with very little effort. There exist many systems for revision control, but their common trait is that they all track changes in some way, and most importantly that these changes can be undone.

4.2.2 Build automation

A different aspect of CI is build automation. Using a timer or trigger, the build automatically runs with the latest updates, preferably from an RC system. This way, the binaries are always up-to-date and the developer does not need to wait for compilation to run the latest tests. Although the latter is less important in a proper CI system, as will be discussed further on.

4.2.3 Test automation

As the code is built at scheduled times, and testing is needed regardless, it makes perfect sense to add a testing step to the automated build. Automating tests saves the developers yet another part of their time, thus freeing more for the actual development and debugging steps. A good CI system can read test results and reports and preferably visualise them in a way that makes them easy to read and locate faults.

Combining all of these practices saves developers, and by extension the company they might work for, a lot of time and money. Considering today's competitive market for software development, any edge that can be obtained is a plus. Even more so when plenty of free continuous integration solutions exist that employ open or widely used standards [34–36].

4.3 xUnit

xUnit is a collection of frameworks that all follow the same basic principle. Its specification defines several key components for a testing framework. They encompass an implementation of a *test case* which contains a single unit test. These test cases are then enveloped in *test suites*. The test suites are groups of tests that need the same conditions before they can be executed. There are more implementation details, but these will not be discussed [37].

JUnit is an implementation for the Java programming language of the xUnit specification. A useful part of JUnit is a standard format for the test reports. The reports are written in eXtensible Markup Language (XML), which is a type of language used for formatting data. The JUnit format is supported by numerous (open source) tools, which makes it an ideal candidate for further investigation [38–40].

4.4 Python

Python is a high-level computer programming language that first appeared in 1991, with the third major revision being released in 2008. It supports object-oriented and structured programming and is easy to use as a scripting language. A great feat of Python's community is that there are many (open source) libraries available for just about any function that comes to mind. This on top of the already impressive amount of libraries that Python itself supports. In addition to this, Python has all of its standard features explained in great detail with good examples in its on-line documenting system. The combination of these features allows any programmer, even with very little know-how, to quickly put together anything that they desire [41].

Part II

Developing a framework

In part I different aspects from both the software and VHDL development worlds were discussed in varying levels of detail. In this part it is explained how bits and pieces of each subject were combined to form a new whole, a VHDL development framework.

5 Outlining

Before work began on developing the framework, there needed to be some boundaries set and goals to work towards. As VHDL development is done on both Windows and Linux, it made sense to keep the whole framework as platform independent as possible. To add to this, the framework needed to be developed in the short span of under a year by one developer who had little experience doing so. The components chosen thus had to be easy to master and work on all platforms with a minimum of modifications.

5.1 First draft

In section 3 some current industry practices were discussed and one of the very basic VHDL testing methods was found to be assertions. In order to remain as broad as possible, work started with processing these assertions first. As mentioned in section 4.1, unit testing is a method that tests the smallest possible part of a code's behaviour. Applied to VHDL, a unit test would then be the part of the code that the assertion checks. This under the assumption that the assertion checks for simple boolean truths such as *output q equals '1'* and not a complex mix of logic.

To the likeness of unit testing, these parts of code were separated into their own test benches. Methods of separation were not available, and thus one was developed. For complete accuracy, a lexical analyzer for VHDL would have been needed. However, the development of such was worth its own thesis and as such a more simple method was used. Instead, estimation of the level of depth in the test bench structure was made using VHDL keywords such as *architecture*, *process* and *begin*.

Furthermore, because lexical analysis was out of the question, the constructs that each assert relied on were unable to be identified. A solution was found in the form of procedures and functions. These were placed in the architecture header, and this way, single-line unit test could be called from the architecture body. All that was left for the parser to do was to find the location of these lines and place each of them in a newly created test bench.

5.2 Draft review

Even though it sounds simple, this method was still needlessly complicated and a better method had to be devised. On top of this, the parser created a heap of little, separate test benches which added more work for the developer as he had to execute all of these separately. A possible solution was the automatic execution and result capture of the created test benches. This would reduce the needed work to a light modification of the test bench and a single execution of the parser.

5.3 Fresh start

As everything up until now was based on code created on the fly, it became clear that the original files weren't ready to be refactored or otherwise modified. The experience gained from writing everything mentioned in section 5 was put to good use, and a clean build was started. With this build a number of assumptions were made:

- Optional arguments to modify the behaviour should be specified
- Full on automation should be included
- A log detailing events should be kept
- Proper documentation should be provided

All of these modifications made the framework more suitable for developers. The optional arguments would allow for a greater control over the program. The automation countered the extra work created from the generation of all the separate test benches. The log made it easier to figure out where things went wrong should they go wrong. And lastly, the documentation was a given for any decent project. In figure 3 there is a block schematic overview of the inner workings of the new framework. It is explained in section 6 how this framework should be used.

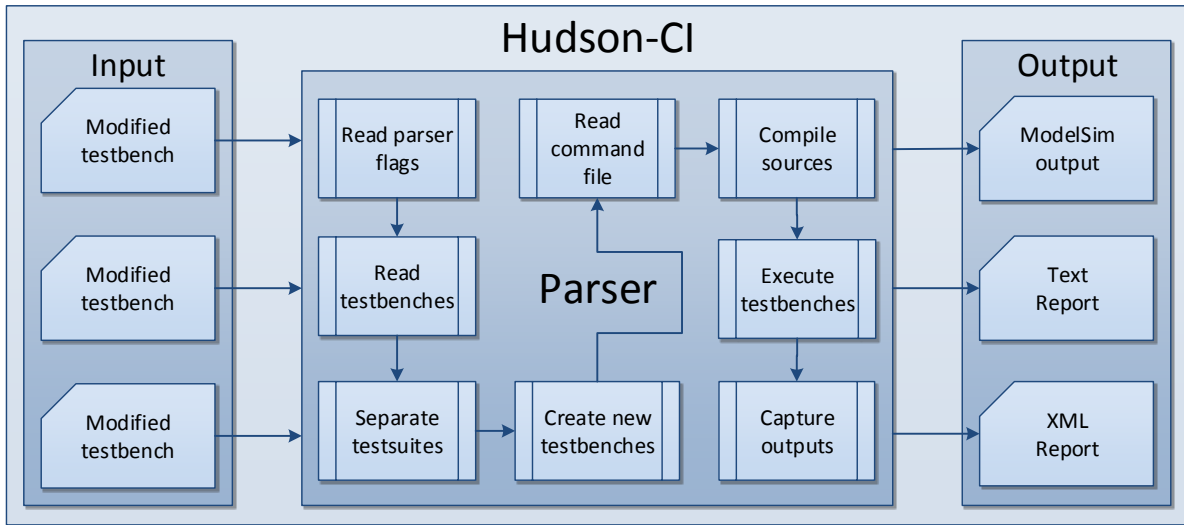


Figure 3: Workflow of the framework

6 Using the framework

For the framework to be properly usable, a number of things had to be considered. Firstly, not all test benches were suited to be processed by the parser; it relied heavily on the independence of tests from one another. Therefore, tests that relied on outputs from previous tests had to be kept together, which lead to a large portion of such test benches needing to be completely rewritten.

Secondly, even with properly suited test benches, the output was to remain uniform and recognizable. As mentioned in section 5.1, to have this done automatically would have required a full blown lexical analysis which was simply too complex. As a solution, the use of a testing library was proposed which would provide uniform tests and output.

6.1 Preparing the test bench

For the test bench to be properly parsed, it needed some preparation first. The blocks of tests, henceforth referred to as test suites, were to be independent and separated by keywords that signify the beginning and end of a test suite. Not only this, but the standard asserts were to be replaced with functions from the library, and the library of course needed to be included. An example using the DFF from section 2.1 (full code in appendix A.1):

```
...
    assert q = '0'
        report "Wrong output value at startup" severity FAILURE;
    d <= '1';
    WAIT FOR clk_period;
    assert q = '1'
        report "Wrong output value at first test" severity FAILURE;
...
```

Would turn into (See appendix A.2):

```
...
--Test 1
check_value(q = '0', FAILURE, "Wrong output value at startup");
write(d, '1', "DFF");
check_value(q = '1', FAILURE, "Wrong output value at first test");
...
--End 1
...
```

In the full code it is shown that the functions are overloads of existing functions in the *BitVis* library, which will be discussed in section 6.5. Also visible here are the keywords *--Test* and *--End* by which the test suites are separated. The number next to it is a test suite identifier, but was not necessary as the parser kept track of the count.

Overloading functions, using libraries and seemingly completely overhauling the test bench may appear to be a lot of work for a small test bench such as this. While true, this method was easily scalable for test benches of much larger sizes and with designs of greater complexity. If the developer created a test bench starting with this approach, no time would have been wasted rewriting any code and the readability would have been greatly improved. Lastly, automated logging was built into these procedures to keep track of progress during test runs. This allowed for easier and uniform report generation (see section 6.3).

6.2 Calling the parser

The parser was created as a script, written in Python, thus it was necessary to have access to a command-line interface as Python requires one to run its scripts. Originally, using the parser was as easy as typing `python testbench_parser.py tb_dff.vhd` in the command-line. There was one test bench to be parsed and everything happened automatically according to default methods, hardcoded in the script. During further development it became clear that there were many ways a developer might want to test their project. As such there was a need for a way to modify the parser's behaviour when it ran. Possible solutions included different scripts or a settings file, but ultimately the decided upon method was the use of arguments.

6.2.1 Arguments to the parser

Python had a built-in argument parser, in the library *argparse*. Using this, it became possible to have a powerful but simple to use command-line interface with both optional and required arguments. Included in the library was the option for a clear and thorough help option, providing users with all the information they needed to run the script. A simple example is listed below:

```
$ python src\testbench_parser.py -m partitioned -l Project\sim\tb_dff_r.vhd -c
```

By using this command, the test bench from appendix A.2 was processed using the *partitioned* method, which is the method that uses the keywords *Test* and *End* (see section 6.2.2). The *-c* flag ensured that output of the simulator (in this case ModelSim) was displayed and not just captured and stored as was default. Below are all the flags and their help text:

- '-c', '--cmd' specifies output to be displayed on the command line
- '-d', '--dest' specifies a custom path for the log
- '-f', '--file' specifies -l/--list is a file with a list of .vhd files
- '-l', '--list' specifies the list of .vhd files to be processed (at least one)
- '-m', '--method' specifies what method was used to write the test bench
- '-p', '--precompiled' specifies location of the precompiled dependencies, requires full path
- '-r', '--' specifies custom arguments for simulation start
- '-v', '--version' specifies a different VHDL version, default is 2008

With this parser, it also became possible to run multiple test benches with one call to the script. This allowed testing multiple levels of a hierarchic design, or even completely different designs in one run. Most of the flags are self-evident, but those that are not are explained below.

The *-p* or *--precompiled* flag points to a text file that contains instructions very similar to those found in a Unix *Makefile*. They are commands that need to be executed before the test bench can be compiled. This would be all of the source files used, in their specific order and any special libraries that need to be created beforehand. A full editor might have recognized the dependencies and libraries on its own and created a Makefile or similar accordingly, but here the lack of a decent lexical analysis prevented this.

The *-r* or *--runops* flag is to specify a custom command for ModelSim to run during the execution of the newly made test benches. The default value is *-do "run -all;exit"*, which simply starts the simulation until its end and then exits. However, developers might want to load different memory files in a memory or specify a custom file with more elaborate commands.

6.2.2 Processing methods

In section 6.2.1 the processing method *partitioned* was mentioned as well as the *-m* or *--method* flag. The parser thus supported multiple methods of dividing the test bench, they are listed below:

- Start/stop with identifiers
- Line by line
- Partitioned with identifiers

In the start/stop method, every single line between certain keywords was assumed to contain a test. It was assumed that everything above the keywords was architecture header and everything below was a general ‘end of the architecture body’. This method was the easiest to parse, it looked for *--Start* and created a new test bench with every line until *--Stop*.

In the line method, the parser looked for a process. This process should contain one test per line. It tried to identify this process on its own and as such, could easier make a mistake than the start/stop method.

Finally, in the partitioned method, the developer indicated which blocks are test suites by using the *--Test* and *--End* keywords. Everything in between these keywords was assumed to be a single block of test cases, to be put in the same new test bench. As such, and unlike the previous two methods, it was possible to define multiple test suites.

6.3 Report and log generation

After the parser from section 6.2 had run its course, results were captured automatically. However, making sense of a whole lot of command line output from a text file was challenging. ModelSim made this more difficult by adding a header and a lot of information that isn’t relevant at that time of development. Information such as which files were loaded prior to the execution and which preference file was being read, were not needed for a quick overview. However, in the event of severe failure they could still prove useful.

6.3.1 Basic report

A great deal of filtering of the output file was required, only the passed and failed test results were needed for a proper report. Whether tests passed or failed was outputted with easily recognizable keywords. A simple filtering on these words per line of the output file returned every test result, neatly ordered in passed and failed groups. A total test count was tallied from the combination, and as such a crude report was written to a text file.

6.3.2 XML report

Previously discussed in section 4.3, the JUnit implementation of xUnit uses XML for report viewing. The editor that was used throughout the development of the thesis is *Eclipse*, which was written in Java, the same language JUnit is an implementation for. They could be read and displayed in Eclipse by manually navigating to them or editing the project properties to automatically read them.

To create the XML reports, a pre-made, open-source Python implementation of a JUnit report generator, aptly named `python-junit-xml`, was used. This package required its own external component to be installed, called `setuptools` [39, 42]. In this thesis, versions 1.0 and 3.5.1 respectively were used, but versions 1.3 and 8.0.4 are available at time of writing.

6.3.3 Logkeeping

Finally, a log was kept detailing events throughout the parser’s execution. It contained critical events along with timestamps for precise debugging. Everything within this log file was related only to the parser and its workings. The file that contained details about ModelSim’s execution was not saved.

6.4 Hudson-CI

For full automation, a CI system was integrated into the framework. The CI solution that was incorporated was Hudson-CI, which provided an extensive range of features, including:

- Timed and triggered building from an RC repository
- Automated testing of builds
- Reading reports in the JUnit format
- Graphical and statistical overview of test results throughout builds

Hudson was available as a running `.war` file (a Java web application that runs on a local machine) or as an installed service [43]. The version used throughout this thesis is 3.2.0, but at time of writing 3.2.1 is available.

6.4.1 Using Hudson-CI

For ease of use, Hudson was installed as a local service. It is worth mentioning that to be able to use ModelSim with Hudson as service, Hudson needs to run under the same account as the one ModelSim has a license for.¹ By default it is accessed at port 8080 of the local machine, usually by typing `localhost:8080` in a browser. The following steps were taken to successfully complete several test runs on a Windows 8 machine:

1. Create a new job with:
 - (a) Git repository as RC system
 - (b) The batch command listed in section 6.2.1.
 - (c) JUnit test results published under `VHDL_TDD_Parser/**/*.xml`
 - (d) Manual build activation
2. Press the *Build now* button

This automatically downloaded the source, compiled the external test bench for the DFF, executed the different architectures, captured and formatted the output and published the XML report. For successive runs only step 2 had to be repeated, the code is automatically updated from git should it change. An example of captured statistics is shown in appendix B.1.

¹On Windows 8: View local services > Hudson > right-click: Properties > Log On (tab)

6.5 Bitvis utility library

Throughout the development of the framework multiple approaches were used to enhance developer experience, including a library with widely used and usable functions. At first this was a self developed library, however, after a short while the existing *Bitvis utility library* was used. It was developed by the Norse company Bitvis for verification and is compatible with all versions of VHDL [44]. The library promoted overloading its functions with implementations for the project that was being developed. It also provided an easy to use logging function with extensive customizability, which made it ideal for use in the framework.

To make use of the library, all that needed to be done was to include it in the project, use its functions and overload where needed. For the framework, the use of the *check* or *check_value* functions was critical. The test parser relied on the output created by these functions to assess test success. An example is listed below (full code in appendix A.2):

```
...
LIBRARY bitvis_util;
USE bitvis_util.types_pkg.ALL;
USE bitvis_util.string_methods_pkg.ALL;
USE bitvis_util.adaptations_pkg.ALL;
USE bitvis_util.methods_pkg.ALL;
...
PROCEDURE log(
    msg    : STRING) IS
BEGIN
    log(ID_SEQUENCER, msg, C_SCOPE);
END;

PROCEDURE write(
    SIGNAL data_target  : IN STD_LOGIC_VECTOR;
    CONSTANT data_value : IN STD_LOGIC_VECTOR;
    CONSTANT msg        : IN STRING) IS
BEGIN
    data_target <= data_value;
    WAIT FOR clk_period;
    log(msg & " write(Target:" & to_string(data_value, HEX, AS_IS, INCL_RADIX)
        & ", " & to_string(data_target, HEX, AS_IS, INCL_RADIX) & ")");
END;
...
    write(d, '1', "DFF");
    check_value(q = '1', FAILURE, "Wrong output value at first test");
...
```

7 Results

A number of open source projects were located at *opencores.org* and used for testing purposes. They were a Secure Hash Algorithm (SHA) 256 bit core, an Advanced Encryption Standard (AES) 256 bit core and a Cyclic Redundancy Check (CRC) core. Each of these projects were supposedly in the final stages of development. The example project from the Bitvis utility library was also used.

Hudson-CI was installed as a service on both a Windows 8 and an Ubuntu 14.04 computer. The compiler used was ModelSim PE 10.3c for Windows and 10.1d for Ubuntu. Both operating systems used versions 1.0 and 3.5.1 respectively of the junit-xml and setuptools Python libraries. For Python itself version 2.7 was used.

Each of the projects had an included test bench that was modified according to the partitioned method mentioned in section 6.2.2. A job was created in Hudson-CI for each project separately with the same parameters mentioned in section 6.4.1. The -m flag was added where needed with the appropriate command file.

Each of the Hudson jobs was run at least 5 times for proper statistics, as well as weeding out non-framework related issues. The automatically captured statistics from section 6.4.1 were reviewed, along with the console output. The main window of Hudson displays job success or failure, *weather* statistics on build stability and test results, gathered from the last 5 runs.

S	W	Job ↓	Last Success	Last Failure	Last Duration	Console	
		VHDL-AES	1 min 15 sec (#30)	3 mo 4 days (#16)	41 sec		
		VHDL-Bitvis	38 sec (#35)	1 mo 7 days (#23)	24 sec		
		VHDL-CRC	3 mo 1 day (#12)	1 min 3 sec (#13)	5,5 sec		
		VHDL-SHA	N/A	52 sec (#5)	6,2 sec		

Figure 4: Hudson jobs status window

From the reports it was concluded that the AES and Bitvis jobs ran with at least one successful run. The results (evident from the weather report) indicated that at least one test for each of the successful jobs had passed. Given that the goal of this thesis was not to develop VHDL, but rather the framework, this could be seen as working. Faults were then deliberately introduced in the Bitvis job to see whether they would be detected. As evident from figure 5 (see appendix B.1), they were identified successfully.

Because the introduced faults only interrupted a small part of the entire test run, the goal of separating tests can be said as having been achieved. Furthermore, with the debugging information given in the Hudson report section, the failed tests were quickly detected out of hundreds of tests.

8 Future work

While a working framework was achieved, many parts are missing for a usable whole. Some of the more critical absent pieces are listed below. Alongside are suggestions that are not necessarily needed for a working framework but might make development more pleasant.

8.1 Wider tool support

In section 3.5 it was decided to support ModelSim as the compiler for the framework. Despite its extensive use in the classroom, it is not the only tool at the top. Different tools might use other methods, accept different input and produce output in a different format. While there are too many tools to support them all, either a general approach can be achieved or the most widely used can be integrated.

On the topic of tools, the current support of ModelSim itself is limited at best. Only the very basic steps have been taken to assure compiling. Complex projects might need an extensive list of options to be run before compiling, during simulation or otherwise.

8.2 Lexical analysis

In section 5 it was decided that lexical analysis was out of the question. However, future modification should include this, as the advantages greatly outweigh the work that it requires. A full lexical analysis should allow for more advanced methods of test bench creation to be supported, as well as a precise integration for current methods. Furthermore, a good analysis could separate a test bench on its own, seeing dependencies and eliminating code duplication. And lastly, a great analysis might even be able to convert test benches to the used methods by recognizing independent groups of test. It could then create its own functions out of them and save developers a lot of time in the process.

8.3 Adapted CI tool

In section 4.2 continuous integration was introduced and in section 6.4 Hudson-CI was proposed as the CI tool. However, Hudson-CI is a software development tool meant for software development practices. Its inner workings have shortcomings that, while not critical, are needed for hardware developers. The JUnit reports, for instance, do not recognize sub-millisecond times, and debugging output of the compiler could be integrated. A dedicated hardware CI tool might be developed to provide full support for all the features that are lacking.

9 Commentaries on developing VHDL

While working on this thesis a great deal of time was spent working with VHDL and test benches in particular. Evaluating, re-evaluating and dissecting these test benches to find ways to improve their workings and their developing environment brought good familiarity with the inner workings. As such, certain aspects of the language and its surroundings or lack thereof were the proverbial thorn in the eye and could stand to be improved.

9.1 An industry stuck in 1993

In section 3 industry standard practices were discussed along with some considerations on the tools used in section 3.5. While work is definitely being done to improve the coding practices, the hardware development world is slow to adapt to any change. The very proof of this is that just now (Q4 2014) slow adaptations are being made for the latest VHDL standard, which was defined in 2008. The lack of willingness to improve upon working methods - "*Don't fix what isn't broken.*" - stifles innovation.

9.2 Lack of reflection or introspection

Reflection and introspection are two programming concepts, the former already existing since assembly languages were used. Reflection means that a programming language can modify parts of its innards at runtime. While this certainly sounds non-synthesizable, it could bring about a revolution in the creation of functions for test benches. Instead of having to overload for an unreasonable amount of different configurations, code would simply adapt to the situation by adjusting its arguments to the right type. This, of course, providing the function's structure continues to work after the modifications.

Introspection can be seen as a part of reflection, it is the ability to examine an object at runtime. While VHDL is not an object-oriented programming language, signals, variables and constants can be seen as a type of object. They have many possible configurations and can have properties such as vector length. Being able to examine these properties at runtime and matching object types against each other would provide for some of the same functionality that reflection gives.

9.3 Hardware developers are not software developers

While considering the previous commentaries, it became clear that a major shortcoming of hardware developers is that they are not trained to develop software. Software developers learn certain practices, methods and group works in the course of their education. Hardware developers usually come from an engineering background where they are taught sciences and knowledge of *what* is being developed. This creates a work environment where seldom there is a developer educated in software development methods. Combined with the slow and stagnant change mentioned in section 9.1, this leads to outdated development practices.

10 Conclusion

At the start of this thesis, it was stipulated that developing VHDL is prone to errors and that the test bench and its environment play an integral part in locating and solving these. As a solution, a testing framework was proposed that would work around a central Python script and integrate ideas from the software development world.

A first draft revealed many shortcomings of the methods investigated and a subsequent revision was scheduled. More detailed outlines provided better planning and easier incorporation of the envisioned methods as well as a better coding practice.

The framework was made to integrate one of the most commonly used simulation tools, ModelSim. Around it, the software development practice of continuous integration was built using the Hudson-CI tool. Hudson-CI was used to retrieve the code from an RC repository, call the script and display the output.

The Python script was made to accept different forms of modified test benches and it was made possible to modify the behaviour of the script with command-line arguments. A full text output report, a text summary and a processed report in the JUnit XML format were made available after compilation and simulation.

Compilation and simulation test runs were performed with a simple example and found to be working as expected. Runs with more complicated test benches were found to be working as intended, but lacking features.

The objective of developing a framework was generally achieved, but it severely lacked in implementation details. Better support for the current tools and support for a larger variety of tools is absolutely necessary for a fully operating framework. Finally, it was most definitely shown that applying software development practices is possible and useful in a hardware development environment.

References

- [1] Roger Lipsett, Erich Marschner and Moe Shahdad. “VHDL - The Language”. In: *IEEE Design & Test of Computers* 3.2 (Apr. 1986), pp. 28–41. DOI: 10.1109/mdt.1986.294900.
- [2] J.H. Aylor, R. Waxman and C. Scarratt. “VHDL - Feature Description And Analysis”. In: *IEEE Design & Test of Computers* 3.2 (Apr. 1986), pp. 17–27. DOI: 10.1109/mdt.1986.294899.
- [3] R.D. Acosta, S.P. Smith and J. Larson. “Mixed-mode simulation of compiled VHDL programs”. In: *1989 IEEE International Conference on Computer-Aided Design. Digest of Technical Papers*. IEEE Comput. Soc. Press, 1989. DOI: 10.1109/iccad.1989.76930.
- [4] T.E. Dillinger et al. “A logic synthesis system for VHDL design descriptions”. In: *1989 IEEE International Conference on Computer-Aided Design. Digest of Technical Papers*. IEEE Comput. Soc. Press, 1989. DOI: 10.1109/iccad.1989.76906.
- [5] R. Camposano, L.F. Saunders and R.M. Tabet. “VHDL as input for high-level synthesis”. In: *IEEE Design & Test of Computers* 8.1 (Mar. 1991), pp. 43–49. DOI: 10.1109/54.75662.
- [6] R. Waxman, L. Saunders and H. Carter. “VHDL links design, test, and maintenance”. In: *IEEE Spectr.* 26.5 (May 1989), pp. 40–44. DOI: 10.1109/6.30762.
- [7] *1164-1993 - IEEE Standard Multivalued Logic System for VHDL Model Interoperability*. 1993. DOI: 10.1109/ieeestd.1993.115571.
- [8] Janick Bergeron. *Writing Testbenches: Functional Verification of HDL Models*. Springer US, 2000.
- [9] D. Biederman. “An overview on writing a VHDL testbench”. In: *Proceedings The Twenty-Ninth Southeastern Symposium on System Theory*. IEEE Comput. Soc. Press, 1997. DOI: 10.1109/ssst.1997.581677.
- [10] F. Wotawa. “Debugging {VHDL} designs using model-based reasoning”. In: *Artificial Intelligence in Engineering* 14.4 (2000), pp. 331–351. ISSN: 0954-1810. DOI: 10.1016/S0954-1810(00)00021-2.
- [11] P. Kission, Hong Ding and A.A. Jerraya. “VHDL based design methodology for hierarchy and component re-use”. In: *Proceedings of EURO-DAC. European Design Automation Conference*. IEEE Comput. Soc. Press, 1995. DOI: 10.1109/eurdac.1995.527446.
- [12] Charles M. Weber, Neil C. Berglund and Patricia Gabella. “Mask Cost and Profitability in Photomask Manufacturing: An Empirical Analysis”. In: *IEEE Transactions on Semiconductor Manufacturing* 19.4 (2006), pp. 465–474.
- [13] R. Reetz, K. Schneider and T. Kropf. “Formal specification in VHDL for hardware verification”. In: *Proceedings Design, Automation and Test in Europe*. IEEE Comput. Soc, 1998. DOI: 10.1109/date.1998.655865.
- [14] W.G. Swavely et al. “A generic VHDL testbench to aid in development of board-level test programs”. In: *Proceedings of AUTOTESTCON '94*. IEEE, 1994. DOI: 10.1109/autest.1994.381506.
- [15] T.W. Williams et al. “Code coverage, what does it mean in terms of quality?” In: *Annual Reliability and Maintainability Symposium. 2001 Proceedings. International Symposium on Product Quality and Integrity (Cat. No.01CH37179)*. IEEE, 2001. DOI: 10.1109/rams.2001.902502.

- [16] Albert Tran, Michael Smith and James Miller. “A Hardware-Assisted Tool for Fast, Full Code Coverage Analysis”. In: *2008 19th International Symposium on Software Reliability Engineering (ISSRE)*. IEEE, Nov. 2008. DOI: 10.1109/issre.2008.22.
- [17] Alfonso Martinez Cruz, Ricardo Barron Fernandez and Heron Molina Lozano. “Automated Functional Coverage for a Digital System Based on a Binary Differential Evolution Algorithm”. In: *2013 BRICS Congress on Computational Intelligence and 11th Brazilian Congress on Computational Intelligence*. IEEE, Sept. 2013. DOI: 10.1109/brics-cci-cbic.2013.26.
- [18] S. Regimbal et al. “Automating functional coverage analysis based on an executable specification”. In: *The 3rd IEEE International Workshop on System-on-Chip for Real-Time Applications, 2003. Proceedings*. IEEE Comput. Soc, 2003. DOI: 10.1109/iwsoc.2003.1213040.
- [19] Ankit Gopani. *Importance of Constrained Random Verification Approach*. 2012. URL: <http://asicwithankit.blogspot.be/2012/04/importance-of-constrain-random.html> (visited on 20/12/2014).
- [20] Nathan Kitchen and Andreas Kuehlmann. “Stimulus generation for constrained random simulation”. In: *2007 IEEE/ACM International Conference on Computer-Aided Design*. IEEE, Nov. 2007. DOI: 10.1109/iccad.2007.4397275.
- [21] D. Deharbe, S. Shankar and E.M. Clarke. “Formal verification of VHDL: the model checker CV”. In: *Proceedings. XI Brazilian Symposium on Integrated Circuit Design (Cat. No.98EX216)*. IEEE Comput. Soc, 1998. DOI: 10.1109/sbcc.1998.715418.
- [22] Aldec Inc. *Active-HDL Configurations*. 2014. URL: https://www.aldec.com/en/products/fpga_simulation/active-hdl (visited on 19/12/2014).
- [23] Cadence. *Incisive Enterprise Simulator*. 2014. URL: http://www.cadence.com/products/fv/enterprise_simulator/pages/default.aspx (visited on 19/12/2014).
- [24] Xilinx. *Vivado Synthesis*. 2014. URL: <http://www.xilinx.com/support/answers/62005.html> (visited on 19/12/2014).
- [25] Altera. *Quartus II support for VHDL 2008*. 2014. URL: http://quartushelp.altera.com/14.1/master.htm#mergedProjects/hdl/vhdl/vhdl_list_2008_vhdl_support.htm?GSA_pos=1&WT.oss_r=1&WT.oss=vhdl%202008 (visited on 19/12/2014).
- [26] Altera. *ModelSim 10.0 release notes*. 2014. URL: http://www.altera.com/download/os-support/release-notes_10.0c.txt (visited on 19/12/2014).
- [27] Martin Fowler. *UnitTest*. 2014. URL: <http://martinfowler.com/bliki/UnitTest.html> (visited on 30/06/2014).
- [28] Scott Wambler. *Introduction to Test Driven Development (TDD)*. 2002. URL: <http://agiledata.org/essays/tdd.html> (visited on 14/03/2014).
- [29] L. Williams, J.C. Sanchez and E.M. Maximilien. “On the Sustained Use of a Test-Driven Development Practice at IBM”. In: *AGILE 2007 Conference (AGILE 2007), 13-17 August 2007, Washington, DC, USA*. AGILE '07. IEEE Computer Society, 2007, pp. 6–14. ISBN: 0-7695-2872-4. DOI: 10.1109/AGILE.2007.43.

- [30] M. Siniaalto and P. Abrahamsson. “A Comparative Case Study on the Impact of Test-Driven Development on Program Design and Test Coverage”. In: *Proceedings of the First International Symposium on Empirical Software Engineering and Measurement. ESEM '07*. IEEE Computer Society, 2007, pp. 275–284. ISBN: 0-7695-2886-4. DOI: 10.1109/ESEM.2007.2.
- [31] B. Thirumalesh and N. Nachiappan. “Evaluating the efficacy of test-driven development: industrial case studies”. In: *ISESE*. ACM, 2006, pp. 356–363. ISBN: 1-59593-218-6. DOI: 10.1145/1159733.1159787. (Visited on 23/01/2007).
- [32] Don Wells. *Code the Unit Test First*. 2000. URL: <http://www.extremeprogramming.org/rules/testfirst.html> (visited on 20/03/2014).
- [33] Hui Liu, Yuan Gao and Zhendong Niu. “An Initial Study on Refactoring Tactics”. In: *2012 IEEE 36th Annual Computer Software and Applications Conference*. IEEE, July 2012. DOI: 10.1109/compsac.2012.31.
- [34] Fazreil Amreen Abdul and Mensely Cheah Siow Fhang. “Implementing Continuous Integration towards rapid application development”. In: *2012 International Conference on Innovation Management and Technology Research*. IEEE, May 2012. DOI: 10.1109/icimtr.2012.6236372.
- [35] James H. Hill et al. “CiCUTS: Combining System Execution Modeling Tools with Continuous Integration Environments”. In: *15th Annual IEEE International Conference and Workshop on the Engineering of Computer Based Systems (ecbs 2008)*. IEEE, Mar. 2008. DOI: 10.1109/ecbs.2008.20.
- [36] Sean Stolberg. “Enabling Agile Testing through Continuous Integration”. In: *2009 Agile Conference*. IEEE, Aug. 2009. DOI: 10.1109/agile.2009.16.
- [37] Martin Fowler. *Xunit*. URL: <http://www.martinfowler.com/bliki/Xunit.html> (visited on 02/06/2014).
- [38] J. Kivi et al. “Extreme programming: a university team design experience”. In: *Canadian Conference on Electrical and Computer Engineering. CCECE '00*. IEEE, 2000, pp. 816–820. ISBN: 0-7803-5957-7. DOI: 10.1109/CCECE.2000.849579.
- [39] Brian Beyer. *Python JUnit XML*. 2014. URL: <https://pypi.python.org/pypi/junit-xml/1.3> (visited on 07/05/2014).
- [40] *JUnit*. URL: <http://JUnit.org> (visited on 25/07/2014).
- [41] *Python*. URL: <https://www.python.org/> (visited on 28/12/2014).
- [42] Jason R. Coombs and Phillip J. Eby. *setuptools*. 2014. URL: <https://pypi.python.org/pypi/setuptools> (visited on 07/05/2014).
- [43] *Hudson-CI*. 2014. URL: <http://hudson-ci.org/> (visited on 25/05/2014).
- [44] *Bitvis Utility Library*. 2014. URL: <http://bitvis.no/products/bitvis-utility-library/> (visited on 25/08/2014).

Appendices

A Code examples

A.1 DFF test bench

```
LIBRARY IEEE;
USE IEEE.std_logic_1164.ALL;

ENTITY tb_dff_r IS
END tb_dff_r;

ARCHITECTURE Behavioural OF tb_dff_r IS
    COMPONENT dff
    PORT(
        d    : IN  STD_LOGIC;
        clk  : IN  STD_LOGIC;
        q    : OUT STD_LOGIC;
    END COMPONENT;

    SIGNAL d    : STD_LOGIC := '0';
    SIGNAL clk  : STD_LOGIC := '0';
    SIGNAL q    : STD_LOGIC := '0';

    CONSTANT clk_period : TIME := 10 ns;

BEGIN
    uut: dff PORT MAP (
        d => d,
        clk => clk,
        q => q
    );
    clk_process : PROCESS
    BEGIN
        clk <= '0';
        WAIT FOR clk_period/2;
        clk <= '1';
        WAIT FOR clk_period - clk_period/2;
    END PROCESS;

    stim_proc: PROCESS
    BEGIN
        WAIT FOR clk_period;
        ASSERT q = '0'
            REPORT "Wrong output value at startup" SEVERITY FAILURE;
        d <= '1';
        WAIT FOR clk_period;
        ASSERT q = '1'
            REPORT "Wrong output value at first test" SEVERITY FAILURE;
        d <= '0';
        WAIT FOR clk_period;
        ASSERT q = '0'
            REPORT "Wrong output value at final test" SEVERITY FAILURE;
        WAIT;
    END PROCESS;
END Behavioural;
```

A.2 Revised DFF test bench

```
LIBRARY IEEE;
USE IEEE.std_logic_1164.ALL;

LIBRARY bitvis_util;
USE bitvis_util.types_pkg.ALL;
USE bitvis_util.string_methods_pkg.ALL;
USE bitvis_util.adaptations_pkg.ALL;
USE bitvis_util.methods_pkg.ALL;

ENTITY tb_dff IS
END tb_dff;

ARCHITECTURE Behavioural OF tb_dff IS
    COMPONENT dff
    PORT(
        d    : IN  STD_LOGIC;
        clk  : IN  STD_LOGIC;
        q    : OUT STD_LOGIC;
    END COMPONENT;

    SIGNAL d    : STD_LOGIC := '0';
    SIGNAL clk  : STD_LOGIC := '0';
    SIGNAL q    : STD_LOGIC := '0';
    CONSTANT clk_period : TIME := 10 ns;

    PROCEDURE log(
        msg    : STRING) IS
    BEGIN
        log(ID_SEQUENCER, msg, C_SCOPE);
    END;

    PROCEDURE write(
        SIGNAL data_target : IN STD_LOGIC_VECTOR;
        CONSTANT data_value : IN STD_LOGIC_VECTOR;
        CONSTANT msg       : IN STRING) IS
    BEGIN
        data_target <= data_value;
        WAIT FOR clk_period;
        log(msg & " write(Target: " & to_string(data_value, HEX, AS_IS, INCL_RADIX)
            & ", " & to_string(data_target, HEX, AS_IS, INCL_RADIX) & ")");
    END;

BEGIN
    uut: dff PORT MAP (
        d => d,
        clk => clk,
        q => q
    );

    clk_process : PROCESS
    BEGIN
        clk <= '0';
        WAIT FOR clk_period/2;
        clk <= '1';
        WAIT FOR clk_period - clk_period/2;
    END PROCESS;

    stim_proc: PROCESS
    BEGIN
        check_value(q = '0', FAILURE, "Wrong output value at startup");
        write(d, '1', "DFF");
        check_value(q = '1', FAILURE, "Wrong output value at first test");
        write(d, '0', "DFF");
        check_value(q = '0', FAILURE, "Wrong output value at final test");
    END PROCESS;

END Behavioural;
```

B Figures

B.1 Hudson-CI example statistics




Figure 5: Hudson-CI example statistics.

