

Building a better VHDL testing environment

Joren Guillaume

Supervisors: Prof. Luc Colman, dhr. Hendrik Eeckhaut (Sigasi)

Counsellor: dhr. Lieven Lemingre (Sigasi)

Master's dissertation submitted in order to obtain the academic degree of
Master of Science in de industriële wetenschappen: elektronica-ICT

Department of Industrial Technology and Construction

Chairman: Prof. Marc Vanhaelst

Faculty of Engineering and Architecture

Academic year 2014-2015



Usage restrictions

The author gives permission to make this master dissertation available for consultation and to copy parts of this master dissertation for personal use. In the case of any other use, the copyright terms have to be respected, in particular with regard to the obligation to state expressly the source when quoting results from this master dissertation.

Preface

Abstract

Using the Document Class phdsymp.cls

Joris Thybaut

Supervisor(s): Eric Laermans, Luc Dupré

Abstract—This article explains how to use the L^AT_EX style recommended for the Proceedings of the FTW PhD Symposium. The article is itself an example of the phdsymp.cls style in action.

Keywords—Style file, L^AT_EX, FTW PhD Symposium

I. INTRODUCTION

THE Symposium covers a wide range of topics and reflects the diversity of research activities at our Faculty, such as:

- Applied Physics
- Architecture
- Automation
- ...

Authors who have prepared their articles using L^AT_EX can get them formatted in the style we would like to recommend for the Proceedings of the Symposium. The style file phdsymp.cls can be used together with the bibliography style file phdsymp.bst.

We recommend a *double column* style as this makes the article easier to read. The column width is 21 pica (approximately 90 mm). In this sample file you will find examples for the layout of displayed equations, theorems, tables, figures, etc.

II. HOW TO USE THE FILE PHDSYMP.CLS

A. General Information

This style file has been written so to allow, with very few changes, the formatting of input that is suitable for the L^AT_EX article style. First, the phdsymp.cls style file has to be selected with a command of the form

```
\documentclass[twocolumn]{phdsymp}
```

The default font size is 10 points.

The Symposium Proceedings will not include author affiliations below or beside the name(s) of the author(s); instead, use the command `\thanks{...}` to list addresses. Note that the `\thanks{. .}` command in the title no longer produce marks: the thanks-footnote should therefore be self-contained, with address and name of the author(s).

The command “`\PARstart{X}{YYY} ZZZ`” produces a large letter X at the beginning of the paragraph. The string YYY will be automatically changed to capital letters.

The bibliography style file phdsymp.bst allows B^IB^T_EX to include the references from the chosen bibliography file(s) according to the format recommended for the Symposium Proceedings.

Footnotes produce a footnote mark as usual.¹

J. Thybaut is with the Chemical Engineering Department, Ghent University (UGent), Gent, Belgium. E-mail: Joris.Thybaut@UGent.be .

¹The footnote is indicated by a footnote mark

In figure ?? we can see an example for the definition of the title page and of the main commands needed to compile a L^AT_EXfile with phdsymp.cls.

```
\documentclass[twocolumn]{phdsymp}

\usepackage{times}

\begin{document}

\title{Using the Document Class phdsymp.cls}

\author{Joris Thybaut
\thanks{J.~Thybaut is...}}

\supervisor{Eric Laermans, Luc Dupr\`e}

\maketitle

\begin{abstract}
This article ...
\end{abstract}

\begin{keywords}
Style file...
\end{keywords}

\section{Introduction}
\PARstart{T}{he} Symposium ...

\bibliographystyle{phdsymp}
\bibliography{bib-file}

\end{document}
```

Fig. 1. Input used to produce this paper.

B. Additional Changes

Most changes resulting from the use of phdsymp.cls should be transparent to the user. For instance, captions for figures and tables have been modified. Caption of tables, however, should be defined before the table item.

B.1 Environments

The environments for theorem, propositions, lemmas, etc. can be defined with the usual L^AT_EX `[?], [?]` command `\newtheorem{...}{...}`. The proof environment is already defined.

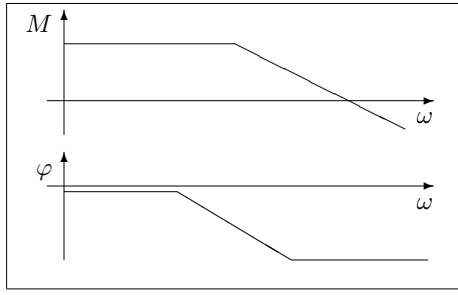


Fig. 2. This is a sample figure. The caption comes after the figure.

TABLE I
THE CAPTION COMES BEFORE THE TABLE.

	title page	odd page	even page
onesided	leftTEXT	leftTEXT	leftTEXT
twosided	leftTEXT	rightTEXT	leftTEXT

Theorem 1 (Theorem name) Consider the system

$$\begin{aligned}\dot{x} &= Ax + Bu \\ y &= Cx + Du\end{aligned}\quad (1)$$

If A is stable, then the pair $\{A, B\}$ is stabilizable. Moreover, this holds for any B .

Proof: The proof is trivial. ■

III. OPTIONAL FORMATTING

When you are happy with the *ultimate unequivocal final version*, you may perform following additional changes, although this is certainly not necessary.

A. “Hard-Coding” Symbolics

Change the symbolics so that the file actually contains the reference numbers *i.e.* “... \cite{fred:88} ...” should be changed to “... [3] ...”. One author (who used the style file) did a smart thing *after* he had decided upon a final version. He put his \cite{...} command and other symbolics on a line on their own and commented them out (from the formatting) by putting a % sign before each symbolic. Then, on the next line he just inserted the copy-matching numerical, like this:

```
Well, according the Fred Bloggs
%\cite{fred:88}
[24]
the value of  $\alpha$  should be even
greater than what we think it should be.
```

Thus, *he* knows he put in the correct (copy-matching) numerical and the publishing staff can send him back an author-proof that correctly matches his submission. This is not so useful here as we do not expect you to send the \LaTeX file, but a PDF-version of your article.

The above also applies to the referencing of table and figures (and any “auto-numbering” feature, standard or synonymous with your system).

Figure captions can be part of the text (in between paragraphs) like this:

And in Fig. 3 we see that the value of α increases exponentially.

Fig. 3\quad This is the caption for figure 3 showing some α .

And after the caption we continue on with the next paragraph, like this.

In essence, by you actually putting in the *correct copy matching* numerals so that no problems arise with incomplete files being sent to the transactions (the wrong *.bib, *.bbl files, the wrong versions of figures etc). Also, and more importantly, the numbers that are on your hard-copy (and in the reviewer’s hands) will be the same ones that you receive in your author proof. Again, this is not so useful here as we do not expect you to send the \LaTeX file, but a PDF-version of your article.

B. Including the Bibliography into the \LaTeX Source File

You can reduce the number of files you have to send to the publishers in the following way. Run \BibTeX on the *.aux file. This creates a *.bbl file: include this into your \LaTeX source file at the place where you defined the \bibliography{...} command and comment this command out. Remove the *.bbl file. Then, your \LaTeX file will include all the necessary information about your bibliography and no *.bbl or *.bib file will be needed.

This may seem like an awful lot of work... but not really.. This will allow to process your paper quickly and efficiently, and assure you that what you send in *will* actually be sent back to you without mistakes (cites, refs etc.).

However, this would only have been useful if you had been requested to send the \LaTeX file itself instead of a PDF-version of the article.

IV. CONCLUSIONS

This sample article has presented the style file phdsymp.cls This file can be especially useful in preparing articles for submission to the FTW PhD Symposium.

ACKNOWLEDGMENTS

The authors would like to acknowledge the suggestions of many people.

REFERENCES

- [1] Leslie Lamport, *A Document Preparation System: \LaTeX , User’s Guide and Reference Manual*, Addison Wesley Publishing Company, 1986.
- [2] Helmut Kopka, *\LaTeX , eine Einführung*, Addison-Wesley, 1989.
- [3] D.K. Knuth, *The \TeX book*, Addison-Wesley, 1989.
- [4] D.E. Knuth, *The METAFONTbook*, Addison Wesley Publishing Company, 1986.

Contents

List of Acronyms	11
I Problem and background	12
1 Problem	12
2 Introduction	13
2.1 Digital Electronics	13
2.2 Hardware Description Languages	14
2.3 Test Driven Development	15
2.3.1 Unit Testing	16
2.3.2 Test First Development	16
2.3.3 Refactoring	16
2.3.4 Test Driven Development	16
3 Testing VHDL	17
4 Building a framework	18
4.1 Continuous Integration	18
4.1.1 Revision Control	18
4.1.2 Build Automation	18
4.1.3 Test Automation	18
4.2 Modelsim	19
5 The future of testing	20
6 Conclusion	21

List of Figures

1 Typical CMOS design flow 15

List of Tables

List of Acronyms

CI	Continuous Integration	18
CMOS	Complementary Metal Oxide Semiconductor	13
DFF	D Flip-Flop	14
DUT	Device Under Test	15
HDL	Hardware Description Language	14, 19
IEEE	Institute of Eletrical and Eletronics Engineers	14
NAND	Not AND	13
NOR	Not OR	13
RC	Revision Control	18
RTL	Register Transfer Level	14
TDD	Test Driven Development	15, 16
TFD	Test First Development	16
UUT	Unit Under Test	15
VHDL	VHSIC Hardware Description Language	12, 14, 18, 19

Part I

Problem and background

1 Problem

Developing VHSIC Hardware Description Language (VHDL), like any code, is prone to error creation, either by user or by wrong product specifications. To ensure errors are weeded out before the more expensive production begins, the code must be subjected to rigorous testing. Currently, large and impractical tests are needed to fully test a product.

Because testing is such a time consuming process, finding errors often results in severe delays due to the need to both correct the error and test for others. Therefore it is in the best interests of both testing- and software engineers to write testbenches with maximal coverage, optimal readability and minimum time spent. It is also important to find and correct errors with minimal delay and maximal efficiency. This entire process should affect the least amount of code possible in order to minimize time spent retesting and modifying the code.

In this thesis, the objective is to provide VHDL users with an operating system independent framework. This framework will allow users to quickly and consistently create, modify and execute testbenches. To accomplish all of this, a number of tried and proven external tools will be gathered around a central python script. This combination will ensure timely and automated building, testing and test report generation.

2 Introduction

2.1 Digital Electronics

There are two kinds of electronic appliances and circuits; digital and analogue. Digital electronics differ from analogue electronics in that they use a discrete set of voltage levels to transmit signals. The most common number of items in the set is 2, a level for one (commonly named *high*) and a level for zero (commonly named *ground* or *low*). The advantage of using a discrete number of levels rather than a continuous signal as is used in analogue electronics, is that noise generated by the environment, thermal noise and other interfering factors will have but a minor influence on the signal.

To process these discrete signals, electronics are made up of transistors that nowadays are formed in with the Complementary Metal Oxide Semiconductor (CMOS) technology. This technology uses both an *NPN* and a *PNP* transistor that work in a push-pull configuration. The p's and n's in NPN and PNP simply stand for *Positive* and *Negative*. They are made of positively and negatively doped lumps of semiconductor, usually silicon-dioxide (SiO_2). A transistor is basically a blockage on a track and depending on the force applied to its *gate*, it opens or closes the track.

In reality, the force takes the form of a current and an NPN transistor opens its gate when a positive current is applied. A PNP transistor, however, always leaves its gate open until a current is applied. This means that if we send the same signal to an NPN and a PNP transistor, with one of the signals inverted, we can open and close two parts of the entire circuit at the same time. This is useful to both direct a certain signal to ground and at the same time close its connection with the *source* (the power source). Hence also the name *Complementary* MOS, the NPN and PNP complement each other.

A certain combination of transistors is used to make *logic gates*. These logic gates make sure that only a certain combination of ones and zeroes at the inputs result in certain ones or zeroes at the outputs. For instance, one of the most common logic gates is a Not AND (NAND) gate. This gate has a number of inputs ranging from 2 to theoretically infinity (but practically 3 or 4) and only outputs a *low* signal if all of the inputs are *high* (digital one), otherwise its output is *at ground* or *low* (digital zero). The other most common logic gate is the Not OR (NOR) gate. This gate outputs a low signal if any of its inputs are high, otherwise it outputs a low.

A common mistake is to think that low or ground mean *zero voltage*. This is only partially true, the high signals are measured with ground as their reference. So a high signal of 1.8 V would be 1.8 V higher than ground, and could be considered to be at 1.8 V if ground is the theoretical zero. These logic gates are themselves combined to build higher-level blocks such as flip-flops, which are used to make registers and so on up to the entire chip design.

2.2 Hardware Description Languages

A Hardware Description Language (HDL) can be used to describe any one of the levels mentioned in the previous section, right down to the logic gate level. However, this last one is wholly unnecessary considering synthesis tools can produce superior logic gate-level layouts [1]. The level that uses certain blocks of logic gates to describe more complex behaviour is called the Register Transfer Level (RTL). Some blocks are standard implementations that have been widely used and are nearly fully optimized, such as memories, flip-flops and clocks. An RTL flip-flop implementation is shown here:

```
library IEEE;
use IEEE.std_logic_1164.all;
entity DFF is
    Port(D    : in std_logic;
         CLK  : in std_logic;
         Q    : out std_logic;
end DFF;
architecture Behavioural of DFF is
begin
    process (CLK)
    begin
        if rising_edge(CLK) then
            Q <= D;
        end if;
    end process;
end Behavioural;
```

The Institute of Electrical and Electronics Engineers (IEEE) library provides a number of extensions on the original VHDL specification that allow a more realistic simulation and description of hardware behaviour. An *entity* defines the inputs and outputs of a certain building block, in this case the D Flip-Flop (DFF). The D stands for Delay, and it simply puts on its output, Q, that which was on the input, D, one clock period earlier. The *architecture*, in this case Behavioural, takes the description of an entity and assigns a real implementation to it. All *processes* are executed in parallel, this does not mean that all are triggered at the same time, nor do they take equally as long to finish, but it means that any process can be executed alongside any other process. In this case there is only one (nameless) process that describes the entire behaviour of the flip-flop. The process waits for the rising edge of the clock, which is a transition from zero to one, after which it schedules the value of D to be put on Q when the next rising edge appears.

This is a basic example of an entity, an architecture and a process. This flip-flop could be used in certain numbers to build a *register*, a collection of ones and zeroes (henceforth referred to as *bits*) that is used to (temporarily) store these values. The register could then be used alongside combinational logic to build an even bigger entity. The idea here is that small building blocks can be combined to produce vast and complex circuits that are nearly impossible to describe in one go. Adding all these layers together also creates a lot of room for error, and having a multi-level design makes it challenging to pinpoint the exact level and location of any errors. Therefore, it is paramount that all code on all levels is tested thoroughly, which is done by use of *testbenches* [2].

Testbenches are made up of code that takes a certain building block, the Unit Under Test (UUT) or Device Under Test (DUT). The testbench then puts a certain sequence of values on the inputs and monitors the outputs. If the device performs normally, the received output sequence should match a certain *golden reference*, the expected output sequence. In these testbenches it is good practice to observe how well a device performs if its inputs behave outside of the normal mode of operation. When all of these tests have finished and the output performs as expected, the device is ready to be put into production or further down the developmental process.

If a device is not tested properly and faults propagate throughout development, they can be very expensive to correct, especially at the stage of production, where a single photomask, used to “print” part of the layout, can easily cost \$100,000 [3] (€81,000 at time of writing). Therefore a large portion of time is spent writing and executing tests[4]. A large number of practices to improve both the speed and quality of testing and coding exist, but the focus chosen in this thesis is Test Driven Development (TDD). This practice has proven to increase test coverage [5], decrease defect density [7] as well as improve code quality [7, 8].

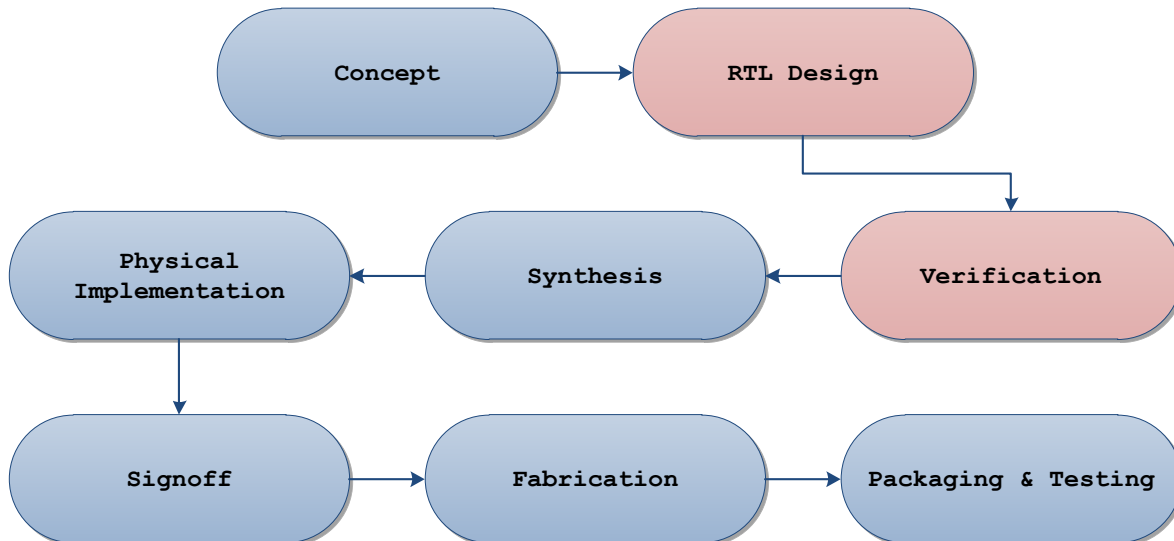


Figure 1: Typical CMOS design flow

2.3 Test Driven Development

Test Driven Development is a proven development technique that has regained traction in the past decade, primarily through the efforts of Kent Beck [?]. The technique focuses on tests being created before the actual code. It is important to make certain distinctions before going more in depth on the used methods. The developing community has a great many practices, each with their own names and methods, and hardly none are mutually exclusive.

2.3.1 Unit Testing

To understand TDD, an basic knowledge of *Unit Testing* is required. In software testing, a unit test is a test designed to check a single unit of code. Ideally, this unit is the smallest piece in which the code can be divided. A unit test should always test only a single entity, and only one aspect of that entity's behaviour. This division in units has a number of benefits, one of the most important being that code is exceedingly easy to maintain. Furthermore, the division of the code makes changes, when needed, fast to be carried out and ensures that only the modified code needs to re-tested.

2.3.2 Test First Development

Another main component of TDD is Test First Development (TFD), a technique that has a developer writing tests first, before any code has been written. This method makes the developer think on what the code has to achieve, rather than what the specific implementation has to be. A key feature of a new test is that it has to fail during its first run. If not, the test is obsolete seeing as the functionality it tests has already been implemented. After being run for the first time (and failing), the developer implements just enough code to get the test to pass. Once the test succeeds, it is time for a new test.

2.3.3 Refactoring

The third pillar of TDD is refactoring. After the tests succeed, it is necessary to clean up the code. A well-done TFD implementation uses the bare amount of code needed to make the test pass, but this is of course usually not the best code possible. Refactoring means that you take existing code and modify only the code itself to perform better. This leaves both the input and output of the tests and code unchanged, only the way the code processes input is altered. It is important in this step to edit nothing in the test code or the outputs or inputs. Otherwise, either the test would behave differently or new tests would have to be written. The latter goes directly against the TFD principle.

2.3.4 Test Driven Development

Test Driven Development is a combination of the previously mentioned techniques. A Unit Test is written before any code, the test is then executed and should fail. After the failure, the most basic code to make the test pass is implemented, the test is then executed again and should pass. After this first pass, the code written for the passing of the test, and only this code, is edited to perform and look better. This follows all steps mentioned above, a *Unit Test* is written according to *Test First Development* and is *Refactored* later.

3 Testing VHDL

4 Building a framework

During the making of this thesis, a number of approaches was investigated and some were put to practical use. Ultimately, the most practical and straightforward approach was to implement a form of Unit Testing. Alongside, an extensive library, written in VHDL, was used to ease the building of testbenches. Finally, a central server was installed that repeatedly and automatically ran the script and read the resulting test outputs.

4.1 Continuous Integration

Continuous Integration (CI) is a software development technique in which developers upload the edits on the software, or their *build*, to a central server which then *continuously integrates* the code from multiple developers. This to prevent integration headaches when the code of multiple developers has diverged to such an extent that it would take much more time to make the edits work together than if they had been integrated early on.

4.1.1 Revision Control

There are many aspects to a properly maintained CI system, but one key aspect of overall programming has to be Revision Control (RC). Not to be confused with the "undo" button in your preferred editor, a good RC system does allow for any and all mistakes made over different edits to be undone with very little work. There exist many systems for revision control, but they all have in common that they track changes one way or the other, and most importantly that these changes can be undone.

4.1.2 Build Automation

A useful but not required aspect is build automation. Using a timer or trigger, the build automatically runs with the latest updates, preferably from an RC system. This way, the binaries are always up to date and the developer does not need to wait for compilation to run the latest tests. Although the latter is less imported in a proper CI system as will be discussed further on.

4.1.3 Test Automation

As the code is build at scheduled times, and testing is needed regardless it makes perfect sense to add a testing step to the automated build. Automating tests saves the developers yet another part of their time, thus freeing more for the actual development and debugging steps. A good CI system can read test reports, or at least some standard of reports.

Combining all of these practices saves developers a lot of time and, by extension, the company they might work for, a lot of money. Considering today's competitive market for software development, any edge that can be obtained is a plus. Even more so when plenty of free Continuous Integration servers exist that employ open or widely used standards. The CI solution that was used throughout this thesis is Hudson-CI. Hudson-CI provides an extensive range of features, including everything listed above. The used features are:

- Timed and triggered building from an RC repository.
- Automated testing of said build.
- Humanly readable reports in the *JUnit* format.
- Graphical and statistical overview of test progress throughout builds.

4.2 Modelsim

As mentioned in section 2 HDLs are used for developing hardware, and need to be tested and build as such. Like any other programming language, they need a dedicated compiler to fault-check the code and build the binaries. Unlike other programming languages, however, they need a simulator in order to verify the builds.

ModelSim is the simulator that was used for several reasons. First, a free student edition was available. Considering licenses outside of school can easily cost upward of \$25,000 this made it ideal for any thesis or student related work. Second, ModelSim supports all versions of VHDL, which is the HDL we are processing. And last, but not least, ModelSim is one of the industry's most used simulators, making the research done of practical use.

5 The future of testing

6 Conclusion

References

- [1] <http://www.asic-world.com/vhdl/intro1.html>
- [2] Writing Testbenches: Functional Verification of HDL Models
- [3] Mask Cost and Profitability in Photomask Manufacturing: An Empirical Analysis
- [4] Citation needed !!
- [5] A comparative case study on the impact of test-driven development on program design and test coverage
- [6] <http://martinfowler.com/bliki/UnitTest.html>
- [7] A Longitudinal Study of the Use of a Test-Driven Development Practice in Industry
- [8] Evaluating the Efficacy of Test-Driven Development

